

2026년도 겨울 인공지능 프로그래밍



위성 격자 자료와 합성곱 신경망(CNN)을 통한 공간 정보 학습

차세대수치예보모델개발사업단

전현주

2026.2.23

실습자료





Hyeon-Ju Jeon Research Scientist at KIAPS

Contact Information

- hjeon@kiaps.org, hgd963@gmail.com
- <https://hgd963.github.io/>


Research Interest

- Multimodal and Irregular Real-world Data Mining
- Applied Machine Learning
- Spatiotemporal Graph Neural Networks
- Application domains: Weather prediction, Disease prediction, Biomedical, Bibliographic network analysis, Anomaly detection, etc.


Professional Experience

- Ph.D., Chung-Ang University (Leave of absence)
- Research Scientist, KIAPS (2021.10 – Present)
- Lecturer, LG CNS (2021 – 2022)

Life Before Deep Learning



Why use features? Why not pixels?




$$f(x_i, W, b) = Wx_i + b$$

Q: What would be a very hard set of classes for a linear classifier to distinguish?

(assuming x = pixels)

Linearly separable classes



Aside: Image Features







Image Features: Motivation



Example: Histogram of Oriented Gradients (HoG)




Divide image into 8x8 pixel regions
Within each region quantize edge direction into 9 bins




Example: 320x240 image gets divided into 40x30 bins; in each bin there are 9 numbers so feature vector has $30*40*9 = 10,800$ numbers

Image features vs ConvNets




Last layer of most CNNs is a linear classifier



Key: perform enough processing so that by the time you get to the end of the network, the classes are linearly separable


Visualizing AlexNet in 2D with t-SNE



(2D visualization using t-SNE)

[Donahue, "DeCAF: DeCAF: A Deep Convolutional ...", arXiv 2013]

Convolutional operator




$$a = [1, 2, 3, 4]$$

$$b = [5, 6, 7, 8]$$

$$a+b = [6, 8, 10, 12]$$


$$a \cdot b = [5, 12, 21, 32]$$

$$a * b = [5, 16, 34, 60, 61, 52, 32]$$




Multiplication

$$[f \cdot g](x)$$




Convolution


$$[f * g](x)$$



Convolution with dice



$$6^2 = 36 \text{ Combinations}$$



Convolution of a_i and b_i

$$(a * b)_n = \sum_{\substack{i,j \\ i+j=n}} a_i \cdot b_j$$

	a_1	a_2	a_3	a_4	a_5	a_6
b_1	$a_1 \cdot b_1$	$a_2 \cdot b_1$	$a_3 \cdot b_1$	$a_4 \cdot b_1$	$a_5 \cdot b_1$	$a_6 \cdot b_1$
b_2	$a_1 \cdot b_2$	$a_2 \cdot b_2$	$a_3 \cdot b_2$	$a_4 \cdot b_2$	$a_5 \cdot b_2$	$a_6 \cdot b_2$
b_3	$a_1 \cdot b_3$	$a_2 \cdot b_3$	$a_3 \cdot b_3$	$a_4 \cdot b_3$	$a_5 \cdot b_3$	$a_6 \cdot b_3$
b_4	$a_1 \cdot b_4$	$a_2 \cdot b_4$	$a_3 \cdot b_4$	$a_4 \cdot b_4$	$a_5 \cdot b_4$	$a_6 \cdot b_4$
b_5	$a_1 \cdot b_5$	$a_2 \cdot b_5$	$a_3 \cdot b_5$	$a_4 \cdot b_5$	$a_5 \cdot b_5$	$a_6 \cdot b_5$
b_6	$a_1 \cdot b_6$	$a_2 \cdot b_6$	$a_3 \cdot b_6$	$a_4 \cdot b_6$	$a_5 \cdot b_6$	$a_6 \cdot b_6$

$$P(\square + \blacksquare = 2) = a_1 \cdot b_1$$

$$P(\square + \blacksquare = 3) = a_1 \cdot b_2 + a_2 \cdot b_1$$

$$P(\square + \blacksquare = 4) = a_1 \cdot b_3 + a_2 \cdot b_2 + a_3 \cdot b_1$$

$$P(\square + \blacksquare = 5) = a_1 \cdot b_4 + a_2 \cdot b_3 + a_3 \cdot b_2 + a_4 \cdot b_1$$

$$P(\square + \blacksquare = 6) = a_1 \cdot b_5 + a_2 \cdot b_4 + a_3 \cdot b_3 + a_4 \cdot b_2 + a_5 \cdot b_1$$

$$P(\square + \blacksquare = 7) = a_1 \cdot b_6 + a_2 \cdot b_5 + a_3 \cdot b_4 + a_4 \cdot b_3 + a_5 \cdot b_2 + a_6 \cdot b_1$$

$$P(\square + \blacksquare = 8) = a_2 \cdot b_6 + a_3 \cdot b_5 + a_4 \cdot b_4 + a_5 \cdot b_3 + a_6 \cdot b_2$$

$$P(\square + \blacksquare = 9) = a_3 \cdot b_6 + a_4 \cdot b_5 + a_5 \cdot b_4 + a_6 \cdot b_3$$


$$P(\square + \blacksquare = 10) = a_4 \cdot b_6 + a_5 \cdot b_5 + a_6 \cdot b_4$$

$$P(\square + \blacksquare = 11) = a_5 \cdot b_6 + a_6 \cdot b_5$$



$$P(\square + \blacksquare = 12) = a_6 \cdot b_6$$

Example: $(1,2,3) * (4,5,6)$

Example: Moving average




Example: 2D




$$\begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$$

$$\frac{1}{9} \begin{bmatrix} 0.5 \\ 0.5 \\ 0.5 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.8 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.8 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.8 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix}$$

Example: 2D



0.003	0.013	0.022	0.013	0.003
0.013	0.060	0.098	0.060	0.013
0.022	0.098	0.162	0.098	0.022
0.013	0.060	0.098	0.060	0.013
0.003	0.013	0.022	0.013	0.003



Example: 2D

kernel

0.25	0.00	-0.25
0.50	0.00	-0.50
0.25	0.00	-0.25



Example: (3,1,4,1,5,9) * (7,7,5)

- Classical signal processing:
 - mean kernel → spatial smoothing
 - Gaussian kernel → physically meaningful smoothing
 - Zero-sum kernel → edge detection
- However, real-world tasks are different:
 - Typhoon structure detection
 - Atmospheric pattern identification
 - Temperature field prediction

What is the correct kernel for a given task?

CNN: Learning the Convolutional operator

- Any linear and shift-invariant operator can be written as a convolution:

$$T(f) = f * w,$$

where w is the kernel.


- In a convolutional neural network:
 - A *convolution layer* implements a linear and shift-invariant operator

$$(a * b)[n] = \sum a[k]b[n - k]$$

- The kernel w is treated as *learnable parameters*
- Training *optimizes the kernel* for the task


Convolutional neural networks

- Layer types:
 - Fully-connected layer
 - Convolutional layer
 - Pooling layer




Fully Connected Layer & Convolution Layer

32x32x3 image -> stretch to 3072 x 1



32x32x3 image



5x5x3 filter



Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

- Fully connected (MLP):

$$y = Wx$$

- Each spatial location has independent weights
- Not shift-invariant and ignores locality




- Convolution:

$$(T * K)(x, y) = \sum_{i,j} T(i, j)K(x - i, y - j)$$

- Same kernel applied at every location (weight sharing)
- Each output is a weighted sum of neighboring values


*For spatially continuous satellite fields,
convolution preserves structures.*

Convolution Layer




Convolution Layer

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:




We stack these up to get a “new image” of size 28x28x6!


Convolution Layer




Convolution Layer




Convolution Layer



What do Conv filters learn?




First-layer conv filters: local image templates
(Often learns oriented edges, opposing colors)



AlexNet: 64 filters, each $3 \times 11 \times 11$

Deeper conv layers: Harder to visualize
Tend to learn larger structures e.g. eyes, letters




6th layer conv layer from an ImageNet model

Visualization from [Springenberg et al, ICLR 2015]

Convolution: Spatial Dimensions

7

7



Input: 7x7
Filter: 3x3
Output: 5x5

Problem: Feature maps shrink with each layer!

In general
Input: W
Filter: K
Output: $W - K + 1$

Convolution: Spatial Dimensions

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Input: 7x7

Filter: 3x3

Output: 5x5

In general

Input: W

Filter: K

Padding: P


Output: $W - K + 1 + 2P$

Problem: Feature maps shrink with each layer!

Solution: Add **padding** around the input before sliding the filter

Receptive fields

Each successive convolution adds $K - 1$ to the receptive field size
With L layers the receptive field size is $1 + L * (K - 1)$



Input


Problem: For large images we need many layers for each output to “see” the whole image

Solution: Downsample inside the network

Output

Slide inspiration: Justin Johnson

Strided Convolution



Input: 7x7
Filter: 3x3
Stride: 2
Output: 3x3

7


In general:
Input: W
Filter: K
Padding: P
Stride: S

Output:
 $(W - K + 2P) / S + 1$

Convolution example

Input volume: $3 \times 32 \times 32$
10 5x5 filters with stride 1, pad 2

Output volume size: ?




Convolution example

Input volume: $3 \times 32 \times 32$

$10 \times 5 \times 5$ filters with stride 1 , pad 2

Output volume size: $10 \times 32 \times 32$

$$32 = (32 + 2 * 2 - 5) / 1 + 1$$




Convolution example

Input volume: **3** x 32 x 32

10 5x5 filters with stride 1, pad 2

Output volume size: 10 x 32 x 32

Number of learnable parameters: ?



Number of learnable parameters: 760

Parameters per filter: **3*5*5 + 1** (for bias) = **76**

10 filters, so total is **10 * 76 = 760**

Convolution example


Input volume: **3** x 32 x 32

10 **5x5** filters with stride 1, pad 2

Output volume size: 10 x 32 x 32

Number of learnable parameters: 760

Number of multiply-add operations?



Number of learnable parameters: 760

Number of multiply-add operations: **768,000**

10*32*32 = 10,240 outputs

Each output is the inner product of two **3x5x5** tensors (75 elems)

Total = $75 * 10240 = \mathbf{768K}$

PyTorch Convolution Layer

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) \[SOURCE\]
```


Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$ can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$


Pooling layer

$64 \times 112 \times 112$




$64 \times 224 \times 224$

pool



224



downsampling

112
112

Given an input $C \times H \times W$,
downsample each $1 \times H \times W$ plane


Hyperparameters:

Kernel Size

Stride

Pooling function


Single depth slice



$64 \times 224 \times 224$




Max pooling with 2x2
kernel size and stride 2



Gives **invariance** to small spatial
shifts. No learnable parameters.


“AlexNet”




“GoogLeNet”



“VGG Net”



“ResNet”




[Krizhevsky et al. NIPS 2012]

[Szegedy et al. CVPR 2015]

[Simonyan & Zisserman,
ICLR 2015]

[He et al. CVPR 2016]

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



High-level Idea: Learn parameters that let us **scale / shift the input data**

1. Normalize input data
2. Scale / shift using learned parameters



Statistics calculated per batch →


Learned parameters applied to each sample →

$$\begin{array}{c} \mathbf{x} : N \times D \\ \boxed{\text{Normalize}} \\ \boldsymbol{\mu}, \boldsymbol{\sigma} : N \times 1 \\ \mathbf{y}, \boldsymbol{\beta} : 1 \times D \\ \mathbf{y} = \boldsymbol{\gamma}(\mathbf{x} - \boldsymbol{\mu}) / \boldsymbol{\sigma} + \boldsymbol{\beta} \end{array}$$

Regularization: Dropout

In each forward pass, randomly set some neurons to zero
Probability of dropping is a hyperparameter; 0.5 is common





Sigmoid


$$\sigma(x) = 1/(1 + e^{-x})$$

- Squashes numbers to range [0,1]
- Historically popular since they have nice interpretation as a saturating “firing rate” of a neuron

Key problem:

Many layers of sigmoids → smaller and smaller gradients.






Q: In which regions does sigmoid have a small gradient?



- Computes $f(x) = \max(0, x)$
- Does not saturate (in +region)
- Very computationally efficient
- Converges much faster than sigmoid in practice (e.g. 6x)

ReLU
(Rectified Linear Unit)

Activation Functions



Case study: VGGNet

Small filters, Deeper networks

8 layers (AlexNet)


-> 16 - 19 layers (VGG16Net)

Only 3x3 CONV stride 1, pad 1
and 2x2 MAX POOL stride 2

11.7% top 5 error in ILSVRC'13


(ZFNet)

-> 7.3% top 5 error in ILSVRC'14



Case study: VGGNet


Q: What is the effective receptive field of three 3x3 conv (stride 1) layers?



Stack of three 3x3 conv (stride 1) layers has same **effective receptive field** as one 7x7 conv layer


But deeper, more non-linearities

And fewer parameters: $3 * (3^2 C^2)$ vs. $7^2 C^2$ for C channels per layer





Case Study: ResNet

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



Case Study: ResNet

What happens when we continue stacking deeper layers on a “plain” convolutional neural network?




56-layer model performs worse on both test and training error
-> The deeper model performs worse, but it's **not caused by overfitting!**

Fact: Deep models have more representation power (more parameters) than shallower models.

Hypothesis: the problem is an *optimization* problem, deeper models are harder to optimize

Case Study: ResNet


Solution: Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping



Case Study: ResNet

Full ResNet architecture:


- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)
Reduce the activation volume by half.



Case Study: ResNet

Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)
Reduce the activation volume by half.



Thank you for listening 😊